

CS162
Operating Systems and
Systems Programming
Lecture 22

Distributed Decision Making (Finished),
TCP/IP Networking, RPC

April 21st, 2020
Prof. John Kubiatowicz
<http://cs162.eecs.Berkeley.edu>

Recall: Distributed Consensus Making

- Consensus problem
 - All nodes propose a value
 - Some nodes might crash and stop responding
 - Eventually, all remaining nodes decide on the same value from set of proposed values
- Distributed Decision Making
 - Choose between “true” and “false”
 - Or Choose between “commit” and “abort”
- Equally important (but often forgotten!): make it durable!
 - How do we make sure that decisions cannot be forgotten?
 - » This is the “D” of “ACID” in a regular database
 - In a global-scale system?
 - » What about erasure coding or massive replication?
 - » Like **BlockChain** applications!

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.2

Recall: Two-Phase Commit

- Since we can't solve the General's Paradox (i.e. simultaneous action), let's solve a related problem
 - Distributed transaction: Two machines agree to do something, or not do it, **atomically**
- Two-Phase Commit protocol:
 - **Prepare Phase:**
 - » The global coordinator requests that all participants will promise to commit or rollback the transaction
 - » Participants record promise in log, then acknowledge
 - » If anyone votes to abort, coordinator writes “Abort” in its log and tells everyone to abort; each records “Abort” in log
 - **Commit Phase:**
 - » After all participants respond that they are prepared, then the coordinator writes “Commit” to its log
 - » Then asks all nodes to commit; they respond with ack
 - » After receive acks, coordinator writes “Got Commit” to log
- **Persistent stable log on each machine:**
 - Help nodes remember what they have said that they would do
 - » If a machine crashes, when it wakes up it first checks its log to recover state of world at time of crash
 - » Log can be used to complete this process such that all machines either commit or don't commit

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.3

Two-Phase Commit: Setup

- One machine (*coordinator*) initiates the protocol
- It asks **every** machine to **vote** on transaction
- Two possible votes:
 - **Commit**
 - **Abort**
- Commit transaction only if unanimous approval

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.4

Two-Phase Commit: Preparing

Agree to Commit

- Machine has **guaranteed** that it will accept transaction
- Must be **recorded in log** so machine will remember this decision if it fails and restarts

Agree to Abort

- Machine has **guaranteed** that it will **never accept** this transaction
- Must be **recorded in log** so machine will remember this decision if it fails and restarts

Two-Phase Commit: Finishing

Commit Transaction

- Coordinator learns *all machines have agreed to commit*
- Record decision to commit in local log
- Apply transaction, inform voters

Abort Transaction

- Coordinator learns *at least one machine has voted to abort*
- Record decision to abort in local log
- Do not apply transaction, inform voters

Two-Phase Commit: Finishing

Commit Transaction

- Coordinator learns *all machines have agreed to commit*
- Record decision to commit in local log
- Apply transaction, inform voters

Abort Transaction

- Coordinator learns *at least one machine has voted to abort*
- Record decision to abort in local log
- Do not apply transaction, inform voters

Because no machine can take back its decision, exactly one of these will happen

Detailed Algorithm

Coordinator Algorithm

Coordinator sends **VOTE-REQ** to all workers


- If receive **VOTE-COMMIT** from all N workers, send **GLOBAL-COMMIT** to all workers
- If don't receive **VOTE-COMMIT** from all N workers, send **GLOBAL-ABORT** to all workers

Worker Algorithm

- Wait for **VOTE-REQ** from coordinator
- If ready, send **VOTE-COMMIT** to coordinator
- If not ready, send **VOTE-ABORT** to coordinator
 - And immediately abort

- If receive **GLOBAL-COMMIT** then commit
- If receive **GLOBAL-ABORT** then abort

Failure Free Example Execution



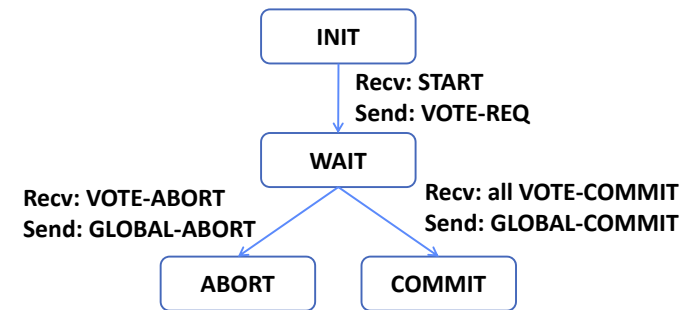
4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.9

State Machine of Coordinator

- Coordinator implements simple state machine:

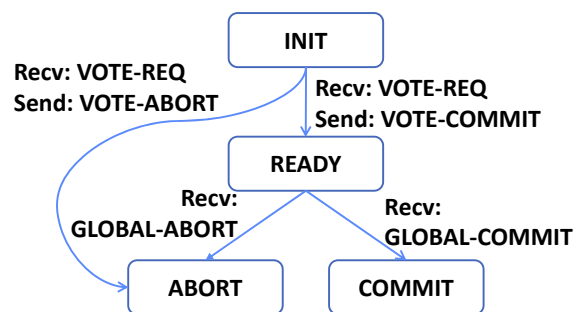


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.10

State Machine of Workers

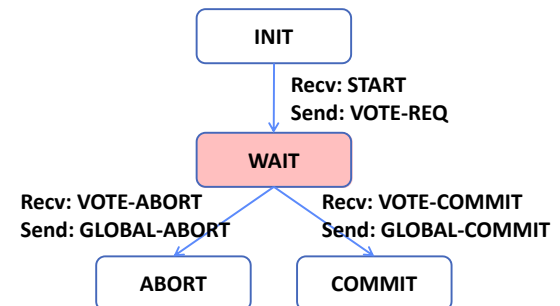


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.11

Dealing with Worker Failures



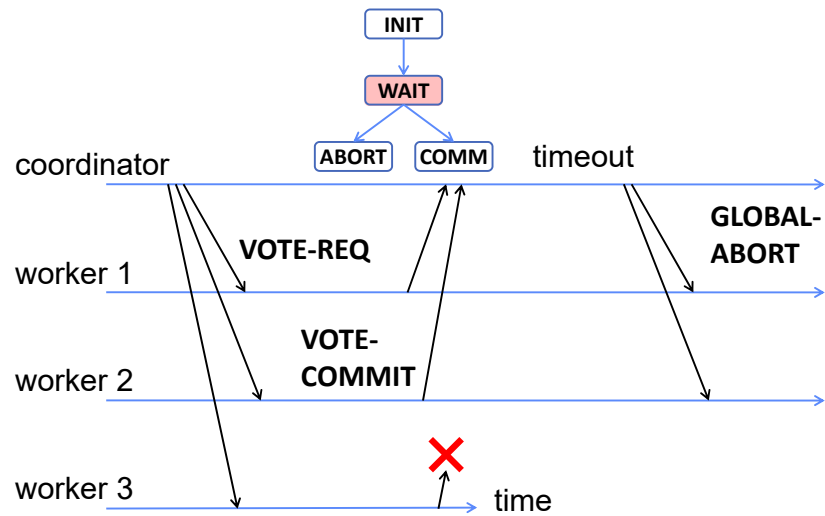
- Failure only affects states in which the coordinator is waiting for messages
- Coordinator only waits for votes in “WAIT” state
- In WAIT, if doesn't receive N votes, it times out and sends GLOBAL-ABORT

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.12

Example of Worker Failure




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.13

Dealing with Coordinator Failure



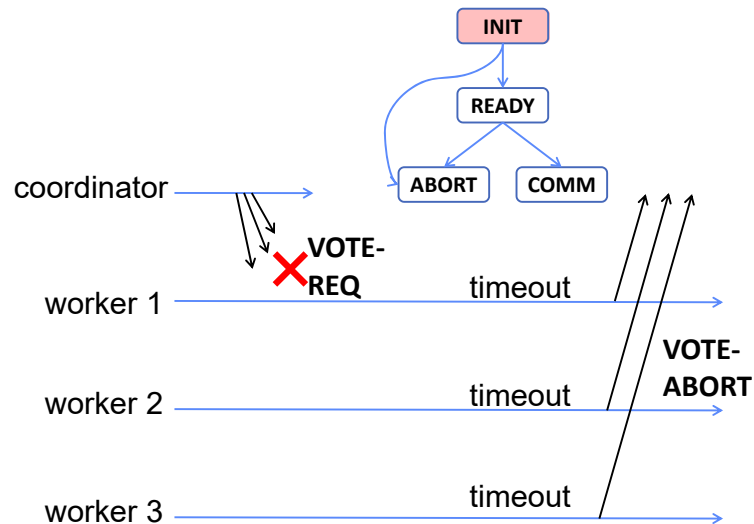
- Worker waits for VOTE-REQ in INIT
 - Worker can time out and abort (coordinator handles it)
- Worker waits for GLOBAL-* message in READY
 - If coordinator fails, workers must **BLOCK** waiting for coordinator to recover and send GLOBAL_* message

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.14

Example of Coordinator Failure #1

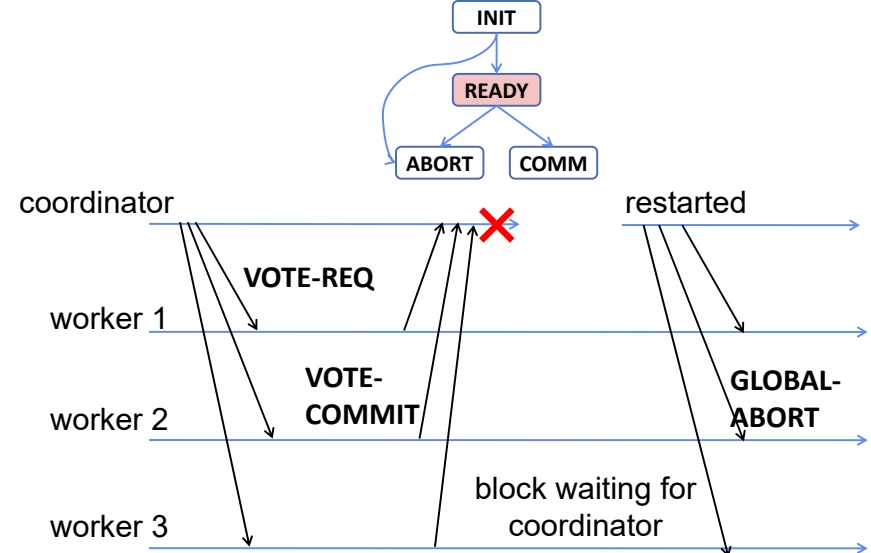


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.15

Example of Coordinator Failure #2



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

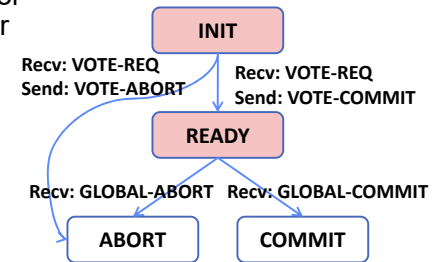
Lec 22.16

Durability

- All nodes use **stable storage** to store current state
 - stable storage is non-volatile storage (e.g. backed by disk) that guarantees atomic writes.
 - E.g.: SSD, NVRAM
- Upon recovery, nodes can restore state and resume:
 - Coordinator **aborts** in INIT, WAIT, or ABORT
 - Coordinator **commits** in COMMIT
 - Worker **aborts** in INIT, ABORT
 - Worker **commits** in COMMIT
 - Worker “**asks**” Coordinator in READY

Blocking for Coordinator to Recover

- A worker waiting for global decision can ask fellow workers about their state
 - If another worker is in ABORT or COMMIT state then coordinator must have sent GLOBAL-*
 - » Thus, worker can safely abort or commit, respectively
- If another worker is still in INIT state then both workers can decide to abort
- If all workers are in ready, need to **BLOCK** (don't know if coordinator wanted to abort or commit)



Distributed Decision Making Discussion (1/2)

- Why is distributed decision making desirable?
 - Fault Tolerance!
 - A group of machines can come to a decision even if one or more of them fail during the process
 - » Simple failure mode called “failstop” (different modes later)
 - After decision made, result recorded in multiple places
- Why is 2PC not subject to the General's paradox?
 - Because 2PC is about *all nodes eventually coming to the same decision – not necessarily at the same time!*
 - Allowing us to reboot and continue allows time for collecting and collating decisions

Distributed Decision Making Discussion (2/2)

- Undesirable feature of Two-Phase Commit: Blocking
 - One machine can be stalled until another site recovers:
 - » Site B writes “prepared to commit” record to its log, sends a “yes” vote to the coordinator (site A) and crashes
 - » Site A crashes
 - » Site B wakes up, check its log, and realizes that it has voted “yes” on the update. It sends a message to site A asking what happened. At this point, B cannot decide to abort, because update may have committed
 - » B is blocked until A comes back
 - A blocked site holds resources (locks on updated items, pages pinned in memory, etc) until learns fate of update

Alternatives to 2PC

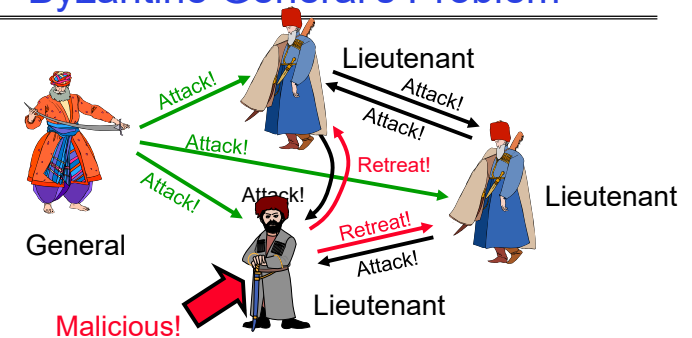
- **Three-Phase Commit:** One more phase, allows nodes to fail or block and still make progress.
- **PAXOS:** An alternative used by Google and others that does not have 2PC blocking problem
 - Develop by Leslie Lamport (Turing Award Winner)
 - No fixed leader, can choose new leader on fly, deal with failure
 - Some think this is extremely complex!
- **RAFT:** PAXOS alternative from John Osterhout (Stanford)
 - Simpler to describe complete protocol
- What happens if one or more of the nodes is malicious?
 - **Malicious:** attempting to compromise the decision making

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.21

Byzantine General's Problem



- Byzantine General's Problem (n players):
 - One General and $n-1$ Lieutenants
 - Some number of these (f) can be insane or malicious
- The commanding general must send an order to his $n-1$ lieutenants such that the following Integrity Constraints apply:
 - IC1: All loyal lieutenants obey the same order
 - IC2: If the commanding general is loyal, then all loyal lieutenants obey the order he sends

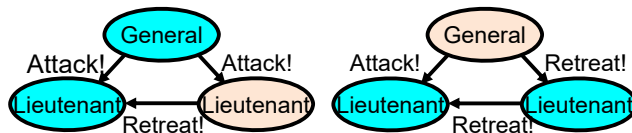
4/21/20

Kubiatowicz CS162 © UCB Spring 2020


Lec 22.22

Byzantine General's Problem (con't)

- Impossibility Results:
 - Cannot solve Byzantine General's Problem with $n=3$ because one malicious player can mess up things



- With f faults, need $n > 3f$ to solve problem
- Various algorithms exist to solve problem
 - Original algorithm has #messages exponential in n
 - Newer algorithms have message complexity $O(n^2)$
 - » One from MIT, for instance (Castro and Liskov, 1999)
- Use of BFT (Byzantine Fault Tolerance) algorithm
 - Allow multiple machines to make a coordinated decision even if some subset of them ($< n/3$) are malicious




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.23

Is a Blockchain a Distributed Decision Making Algorithm?



- Blockchain: a chain of blocks connected by hashes to root block
 - The Hash Pointers are unforgeable (assumption)
 - The Chain has no branches except perhaps for heads
 - Blocks are considered "authentic" part of chain when they have authenticity info in them
- How is the head chosen?
 - Some consensus algorithm
 - In many Blockchain algorithms (e.g. BitCoin, Ethereum), the head is chosen by solving hard problem
 - » This is the job of "miners" who try to find "nonce" info that makes hash over block have specified number of zero bits in it
 - » The result is a "Proof of Work" (POW)
 - » Selected blocks above (green) have POW in them and can be included in chains

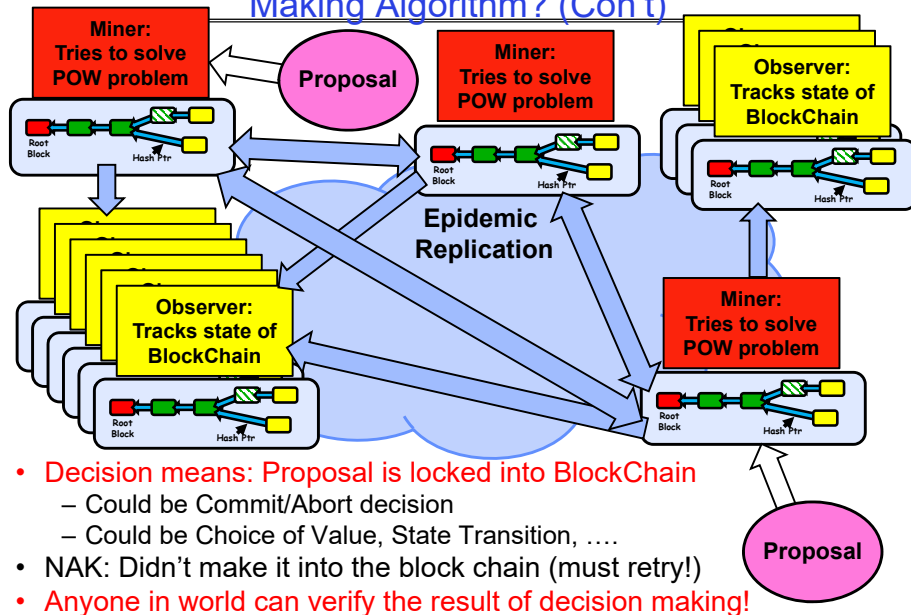
– Longest chain wins

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.24

Is a Blockchain a Distributed Decision Making Algorithm? (Con't)




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.25

Network Protocols

- Networking protocols: many levels
 - Physical level: mechanical and electrical network (e.g., how are 0 and 1 represented)
 - Link level: packet formats/error control (for instance, the CSMA/CD protocol)
 - Network level: network routing, addressing
 - Transport Level: reliable message delivery
- Protocols on today's Internet:



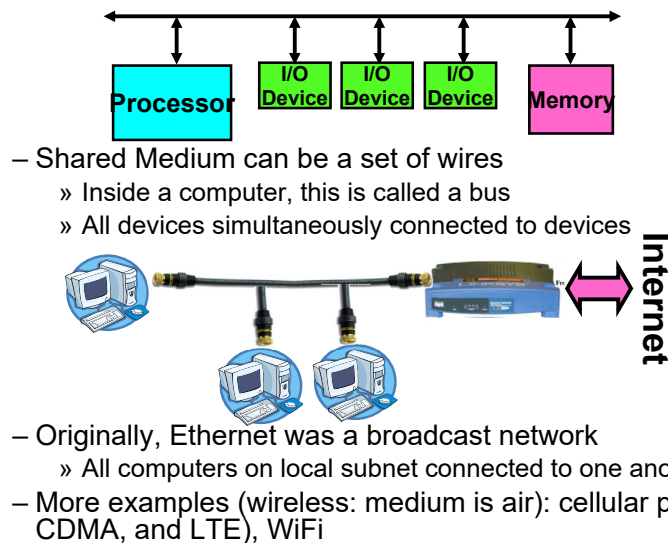
4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.26

Broadcast Networks

- **Broadcast Network: Shared Communication Medium**

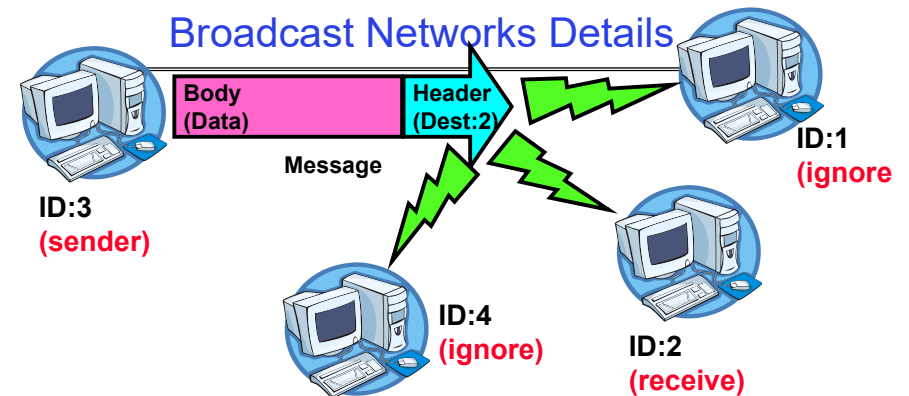


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.27

Broadcast Networks Details



- **Media Access Control (MAC) Address:**
 - 48-bit physical address for hardware interface
 - Every device (in the world!?) has a unique address
- **Delivery:** When you broadcast a packet, how does a receiver know who it is for? (packet goes to everyone!)
 - Put header on front of packet: [Destination MAC Addr | Packet]
 - Everyone gets packet, discards if not the target
 - In Ethernet, this check is done in hardware
 - » No OS interrupt if not for particular destination

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.28

Carrier Sense, Multiple Access/Collision Detection

- Ethernet (early 80's): first practical local area network
 - It is the most common LAN for UNIX, PC, and Mac
 - Use wire instead of radio, but still broadcast medium
- Key advance was in arbitration called CSMA/CD: Carrier sense, multiple access/collision detection
 - Carrier Sense**: don't send unless idle
 - Don't mess up communications already in process
 - Collision Detect**: sender checks if packet trampled.
 - If so, abort, wait, and retry.
 - Backoff Scheme**: Choose wait time before trying again
- How long to wait after trying to send and failing?
 - What if everyone waits the same length of time? Then, they all collide again at some time!
 - Must find way to break up shared behavior with nothing more than shared communication channel
- Adaptive randomized waiting strategy:
 - Adaptive and Random**: First time, pick random wait time with some initial mean. If collide again, pick random value from bigger mean wait time. Etc.
 - Randomness is important to decouple colliding senders
 - Scheme figures out how many people are trying to send!

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.29

MAC Address: Unique Physical Address of Interface

- Can easily find MAC addr. on your machine/device:
 - E.g., ifconfig (Linux, Mac OS X), ipconfig (Windows)


The image shows two screenshots. The left one is a Mac OS 'About' window for a MacBook, showing the Wi-Fi address as 00:1E:C2:CE:12:C4. The right one is a Windows command prompt running 'ipconfig', showing the Physical Address for the Wireless LAN adapter as 00-13-00-E1-11-11 and for the Ethernet adapter as 00-00-00-1A-1F-25. Yellow boxes and arrows highlight these MAC addresses.

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.30

Point-to-point networks



- Why have a shared bus at all? Why not simplify and only have point-to-point links + routers/switches?
 - Originally wasn't cost-effective, now hardware is cheap!
- Point-to-point network**: a network in which every physical wire is connected to only two computers
- Switch**: a bridge that transforms a shared-bus (broadcast) configuration into a point-to-point network
 - Adaptively figures out which ports have which MAC addresses
- Router**: a device that acts as a junction between physical networks to transfer data packets among them
 - Routes between switching domains using (for instance) IP addresses

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.31

The Internet Protocol (IP)

- Internet Protocol: Internet's network layer
- Service it provides: "Best-Effort" Packet Delivery
 - Tries it's "best" to deliver packet to its destination
 - Packets may be lost
 - Packets may be corrupted
 - Packets may be delivered out of order
- IP Is a Datagram service!
 - Routes across many physical switching domains (subnets)



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.32

IPv4 Address Space

- **IP Address:** a 32-bit integer used as destination of IP packet
 - Often written as four dot-separated integers, with each integer from 0–255 (thus representing $8 \times 4 = 32$ bits)
 - Example CS file server is: 169.229.60.83 $\equiv 0xA9E53C53$
- **Internet Host:** a computer connected to the Internet
 - Host has one or more IP addresses used for routing
 - » Some of these may be private and unavailable for routing
 - Not every computer has a unique IP address
 - » Groups of machines may share a single IP address
 - » In this case, machines have private addresses behind a “Network Address Translation” (NAT) gateway
- **Subnet:** network connecting hosts with related IP addresses
 - A subnet is identified by 32-bit value, with the bits which differ set to zero, followed by a slash and a mask
 - » Example: 128.32.131.0/24 designates a subnet in which all the addresses look like 128.32.131.XX
 - » Same subnet: 128.32.131.0/255.255.255.0
 - **Mask:** The number of matching prefix bits
 - » Expressed as a single value (e.g., 24) or a set of ones in a 32-bit value (e.g., 255.255.255.0)
 - Often routing *within* subnet is by MAC address (smart switches)

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.33

Address Ranges in IPv4

- IP address space divided into prefix-delimited ranges:
 - Class A: NN.0.0.0/8
 - » NN is 1–126 (126 of these networks)
 - » 16,777,214 IP addresses per network
 - » 10.xx.yy.zz is private
 - » 127.xx.yy.zz is loopback
 - Class B: NN.MM.0.0/16
 - » NN is 128–191, MM is 0–255 (16,384 of these networks)
 - » 65,534 IP addresses per network
 - » 172.[16–31].xx.yy are private
 - Class C: NN.MM.LL.0/24
 - » NN is 192–223, MM and LL 0–255 (2,097,151 of these networks)
 - » 254 IP addresses per networks
 - » 192.168.xx.yy are private
- Address ranges are often owned by organizations
 - Can be further divided into subnets

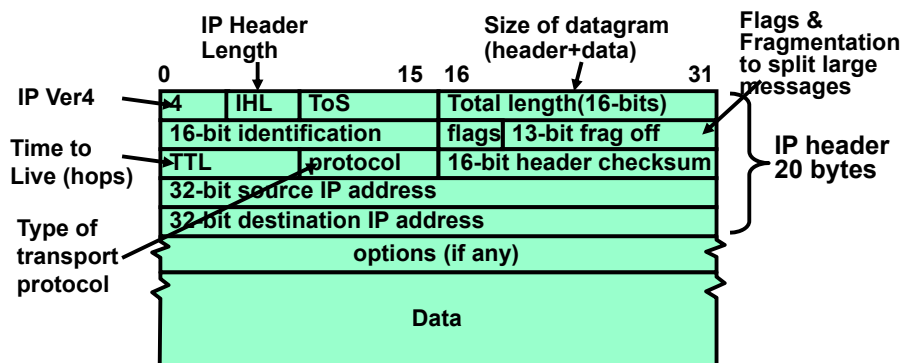
4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.34

IPv4 Packet Format

- IP Packet Format:



- **IP Datagram:** an unreliable, unordered, packet sent from source to destination
 - Function of network – deliver datagrams!


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.35

Wide Area Network

- **Wide Area Network (WAN):** network that covers a broad area (e.g., city, state, country, entire world)
 - E.g., Internet is a WAN
- WAN connects multiple physical (datalink) layer networks (LANs)
- Datalink layer networks are connected by **routers**
 - Different LANs can use different communication technology (e.g., wireless, cellular, optics, wired)




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.36

Routers

- **Forward** each packet received on an **incoming link** to an **outgoing link** based on packet's destination IP address (towards its destination)
- **Store & forward**: packets are buffered before being forwarded
- **Forwarding table**: mapping between IP address and the output link




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.37

Packet Forwarding

- Upon receiving a packet, a router
 - read the IP destination address of the packet
 - consults its forwarding table → output port
 - forwards packet to corresponding output port
- **Default route** (for subnets without explicit entries)
 - Forward to more authoritative router



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.38

IP Addresses vs. MAC Addresses

- Why not use MAC addresses for routing?
 - Doesn't scale
- Analogy
 - MAC address → SSN
 - IP address → (unreadable) home address
- MAC address: uniquely associated to the device for the entire lifetime of the device
- IP address: changes as the device location changes
 - Your notebook IP address at school is different from home



4/21/20

Lec 22.39

IP Addresses vs. MAC Addresses

- Why does packet forwarding using IP addr. scale?
- Because IP addresses can be aggregated
 - E.g., all IP addresses at UC Berkeley start with **0xA9E5**, i.e., any address of form 0xA9E5**** belongs to Berkeley
 - Thus, a router in NY needs to keep a **single** entry for **all** hosts at Berkeley
 - If we were using MAC addresses the NY router would need to maintain **an entry for every** Berkeley host!!
- Analogy:
 - Give this letter to person with SSN: 123-45-6789 vs.
 - Give this letter to "John Smith, 123 First Street, LA, US"



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.40

Setting up Routing Tables


- How do you set up routing tables?
 - Internet has no centralized state!
 - » No single machine knows entire topology
 - » Topology constantly changing (faults, reconfiguration, etc.)
 - Need dynamic algorithm that acquires routing tables
 - » Ideally, have one entry per subnet or portion of address
 - » Could have “default” routes that send packets for unknown subnets to a different router that has more information
- Possible algorithm for acquiring routing table
 - Routing table has “cost” for each entry
 - » Includes number of hops to destination, congestion, etc.
 - » Entries for unknown subnets have infinite cost
 - Neighbors periodically exchange routing tables
 - » If neighbor knows cheaper route to a subnet, replace your entry with neighbors entry (+1 for hop to neighbor)
- In reality:
 - Internet has networks of many different scales
 - Different algorithms run at different scales

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.41

Naming in the Internet

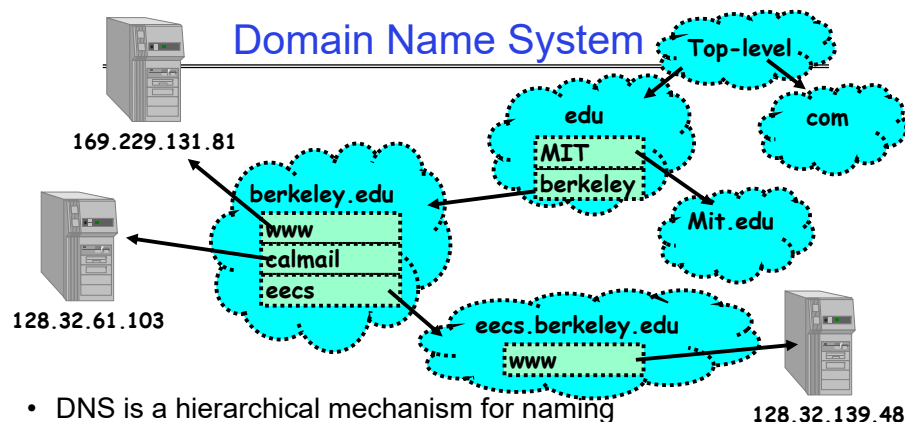


- How to map human-readable names to IP addresses?
 - E.g. `www.berkeley.edu` \Rightarrow `128.32.139.48`
 - E.g. `www.google.com` \Rightarrow different addresses depending on location, and load
- Why is this necessary?
 - IP addresses are hard to remember
 - IP addresses change:
 - » Say, Server 1 crashes gets replaced by Server 2
 - » Or – `google.com` handled by different servers
- Mechanism: Domain Naming System (DNS)

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.42



- DNS is a hierarchical mechanism for naming
 - Name divided in domains, right to left: `www.eecs.berkeley.edu`
- Each domain owned by a particular organization
 - Top level handled by ICANN (Internet Corporation for Assigned Numbers and Names)
 - Subsequent levels owned by organizations
- Resolution: series of queries to successive servers
- Caching: queries take time, so results cached for period of time

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.43

How Important is Correct Resolution?

- If attacker manages to give incorrect mapping:
 - Can get someone to route to server, thinking that they are routing to a different server
 - » Get them to log into “bank” – give up username and password
- Is DNS Secure?
 - Definitely a weak link
 - » What if “response” returned from different server than original query?
 - » Get person to use incorrect IP address!
 - Attempt to avoid substitution attacks:
 - » Query includes random number which must be returned
- In July 2008, hole in DNS security located!
 - Dan Kaminsky (security researcher) discovered an attack that broke DNS globally
 - » One person in an ISP convinced to load particular web page, then all users of that ISP end up pointing at wrong address
 - High profile, highly advertised need for patching DNS
 - » Big press release, lots of mystery
 - » Security researchers told no speculation until patches applied

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.44

Network Layering

- **Layering**: building complex services from simpler ones
 - Each layer provides services needed by higher layers by utilizing services provided by lower layers
- The physical/link layer is pretty limited
 - Packets are of limited size (called the “**Maximum Transfer Unit** or MTU: often 200-1500 bytes in size)
 - Routing is limited to within a physical link (wire) or perhaps through a switch
- Our goal in the following is to show how to construct a secure, ordered, message service routed to anywhere:


Physical Reality: Packets	Abstraction: Messages
Limited Size	Arbitrary Size
Unordered (sometimes)	Ordered
Unreliable	Reliable
Machine-to-machine	Process-to-process
Only on local area net	Routed anywhere
Asynchronous	Synchronous
Insecure	Secure

4/21/20

Lec 22.45

Building a messaging service on IP

- Process to process communication
 - Basic routing gets packets from machine→machine
 - What we really want is routing from process→process
 - » Add “**ports**”, which are 16-bit identifiers
 - » A communication channel (**connection**) defined by 5 items: [source addr, source port, dest addr, dest port, protocol]
- For example: The Unreliable Datagram Protocol (UDP)
 - Layered on top of basic IP (**IP Protocol 17**)
 - » **Datagram**: an unreliable, unordered, packet sent from source user → dest user (Call it UDP/IP)



- Important aspect: low overhead!
 - » Often used for high-bandwidth video streams
 - » Many uses of UDP considered “anti-social” – none of the “well-behaved” aspects of (say) TCP/IP

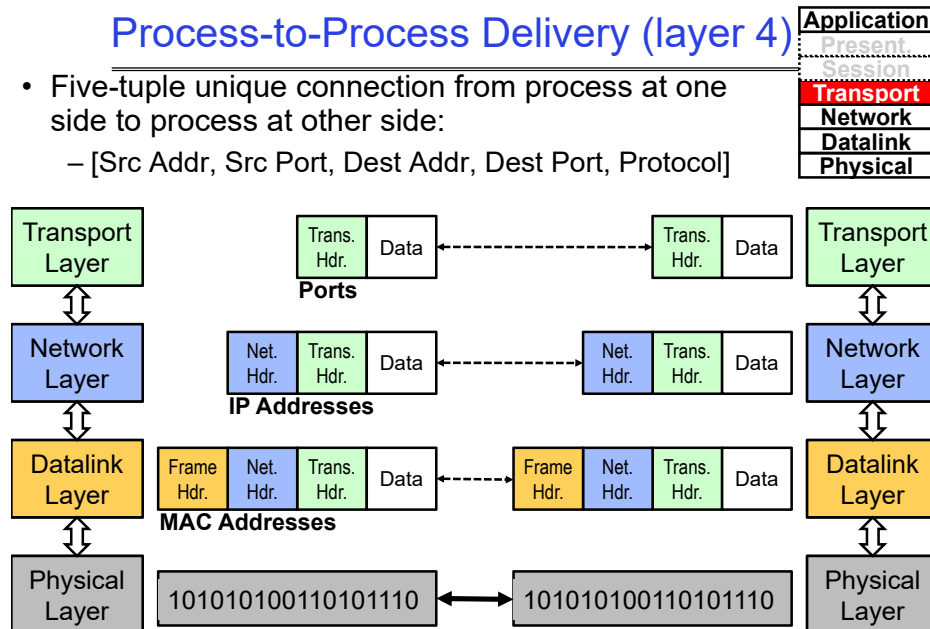
4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.46

Process-to-Process Delivery (layer 4)

- Five-tuple unique connection from process at one side to process at other side:
 - [Src Addr, Src Port, Dest Addr, Dest Port, Protocol]




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.47

Internet Transport Protocols

- Datagram service (**UDP**): IP Protocol 17
 - No-frills extension of “best-effort” IP
 - Multiplexing/Demultiplexing among processes
- Reliable, in-order delivery (**TCP**): IP Protocol 6
 - Connection set-up & tear-down
 - Discarding corrupted packets (segments)
 - Retransmission of lost packets (segments)
 - Flow control
 - Congestion control
- Other examples:
 - DCCP (33), Datagram Congestion Control Protocol
 - RDP (26), Reliable Data Protocol
 - SCTP (132), Stream Control Transmission Protocol
- Services **not available**
 - Delay and/or bandwidth guarantees
 - Sessions that survive change-of-IP-address
 - Security/denial of service resilience/...



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.48

Reliable Message Delivery: the Problem


- All physical networks can garble and/or drop packets
 - Physical media: packet not transmitted/received
 - If transmit close to maximum rate, get more throughput – even if some packets get lost
 - If transmit at lowest voltage such that error correction just starts correcting errors, get best power/bit
 - Congestion: no place to put incoming packet
 - Point-to-point network: insufficient queue at switch/router
 - Broadcast link: two host try to use same link
 - In any network: insufficient buffer space at destination
 - Rate mismatch: what if sender send faster than receiver can process?
- Reliable Message Delivery on top of Unreliable Packets
 - Need some way to make sure that packets actually make it to receiver
 - Every packet received at least once
 - Every packet received at most once
 - Can combine with ordering: every packet received by process at destination exactly once and in order

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.49

Using Acknowledgements



- How to ensure transmission of packets?
 - Detect garbling at receiver via checksum, discard if bad
 - Receiver acknowledges (by sending “ACK”) when packet received properly at destination
 - Timeout at sender: if no ACK, retransmit
- Some questions:
 - If the sender doesn't get an ACK, does that mean the receiver didn't get the original message?
 - No
 - What if ACK gets dropped? Or if message gets delayed?
 - Sender doesn't get ACK, retransmits, Receiver gets message twice, ACK each


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.50

How to Deal with Message Duplication?

- Solution: put sequence number in message to identify re-transmitted packets
 - Receiver checks for duplicate number's; Discard if detected
- Requirements:
 - Sender keeps copy of unACK'd messages
 - Easy: only need to buffer messages
 - Receiver tracks possible duplicate messages
 - Hard: when ok to forget about received message?
- Alternating-bit protocol:**
 - Send one message at a time; don't send next message until ACK received
 - Sender keeps last message; receiver tracks sequence number of last message received
- Pros: simple, small overhead
- Con: Poor performance
 - Wire can hold multiple messages; want to fill up at (wire latency \times throughput)
- Con: doesn't work if network can delay or duplicate messages arbitrarily




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.51

Better Messaging: Window-based Acknowledgements

- Windowing protocol (not quite TCP):**
 - Send up to N packets without ack
 - Allows pipelining of packets
 - Window size (N) < queue at destination
 - Each packet has sequence number
 - Receiver acknowledges each packet
 - ACK says “received all packets up to sequence number X”/send more
- ACKs serve dual purpose:
 - Reliability: Confirming packet received
 - Ordering: Packets can be reordered at destination
- What if packet gets garbled/dropped?
 - Sender will timeout waiting for ACK packet
 - Resend missing packets \Rightarrow Receiver gets packets out of order!
 - Should receiver discard packets that arrive out of order?
 - Simple, but poor performance
 - Alternative: Keep copy until sender fills in missing pieces?
 - Reduces # of retransmits, but more complex
- What if ACK gets garbled/dropped?
 - Timeout and resend just the un-acknowledged packets




4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.52

Transmission Control Protocol (TCP)



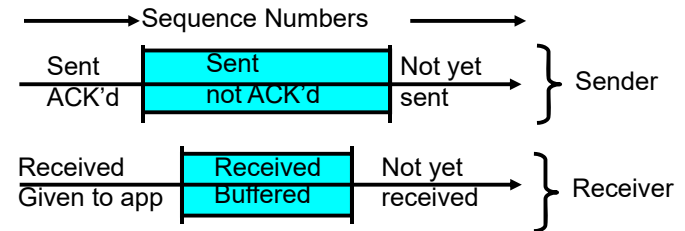
- Transmission Control Protocol (TCP)
 - TCP (IP Protocol 6) layered on top of IP
 - Reliable byte stream between two processes on different machines over Internet (read, write, flush)
- TCP Details
 - Fragments byte stream into packets, hands packets to IP
 - » IP may also fragment by itself
 - Uses window-based acknowledgement protocol (to minimize state at sender and receiver)
 - » “Window” reflects storage at receiver – sender shouldn’t overrun receiver’s buffer space
 - » Also, window should reflect speed/capacity of network – sender shouldn’t overload network
 - Automatically retransmits lost packets
 - Adjusts rate of transmission to avoid congestion
 - » A “good citizen”

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.53

TCP Windows and Sequence Numbers




- Sender has three regions:
 - Sequence regions
 - » sent and ACK'd
 - » sent and not ACK'd
 - » not yet sent
 - Window (colored region) adjusted by sender
- Receiver has three regions:
 - Sequence regions
 - » received and ACK'd (given to application)
 - » received and buffered
 - » not yet received (or discarded because out of order)

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.54

Window-Based Acknowledgements (TCP)



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.55

Congestion Avoidance

- Congestion
 - How long should timeout be for re-sending messages?
 - » Too long → wastes time if message lost
 - » Too short → retransmit even though ACK will arrive shortly
 - Stability problem: more congestion ⇒ ACK is delayed ⇒ unnecessary timeout ⇒ more traffic ⇒ more congestion
 - » Closely related to window size at sender: too big means putting too much data into network
- How does the sender’s window size get chosen?
 - Must be less than receiver’s advertised buffer size
 - Try to match the rate of sending packets with the rate that the slowest link can accommodate
 - Sender uses an adaptive algorithm to decide size of N
 - » Goal: fill network between sender and receiver
 - » Basic technique: slowly increase size of window until acknowledgements start being delayed/lost
- TCP solution: “slow start” (start sending slowly)
 - If no timeout, slowly increase window size (throughput) by 1 for each ACK received
 - Timeout ⇒ congestion, so cut window size in half
 - “Additive Increase, Multiplicative Decrease”

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.56

Open Connection: 3-Way Handshaking

- Goal: agree on a set of parameters, i.e., the start sequence number for each side
 - Starting sequence number (first byte in stream)
 - Must be unique!
 - » If it is possible to predict sequence numbers, might be possible for attacker to hijack TCP connection
- Some ways of choosing an initial sequence number:
 - Time to live: each packet has a deadline.
 - » If not delivered in X seconds, then is dropped
 - » Thus, can re-use sequence numbers if wait for all packets in flight to be delivered or to expire
 - Epoch #: uniquely identifies *which* set of sequence numbers are currently being used
 - » Epoch # stored on disk, Put in every message
 - » Epoch # incremented on crash and/or when run out of sequence #
 - Pseudo-random increment to previous sequence number
 - » Used by several protocol implementations


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.57

Open Connection: 3-Way Handshaking

- Server waits for new connection calling **listen()**
- Sender call **connect()** passing socket which contains server's IP address and port number
 - OS sends a special packet (SYN) containing a proposal for first sequence number, x



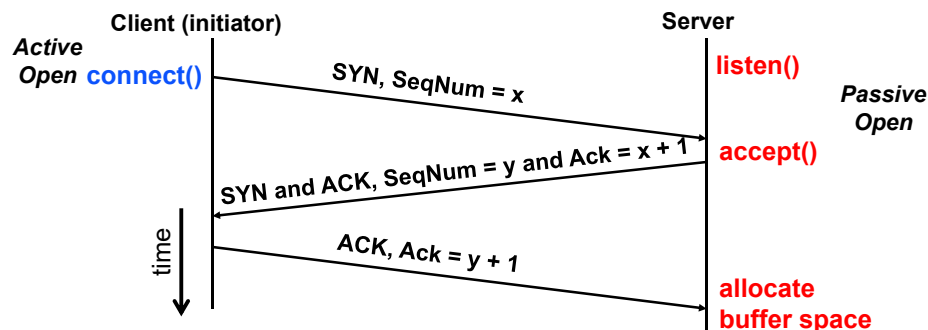
4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.58

Open Connection: 3-Way Handshaking

- If it has enough resources, server calls **accept()** to accept connection, and sends back a SYN ACK packet containing
 - Client's sequence number incremented by one, $(x + 1)$
 - » Why is this needed?
 - A sequence number proposal, y, for first byte server will send

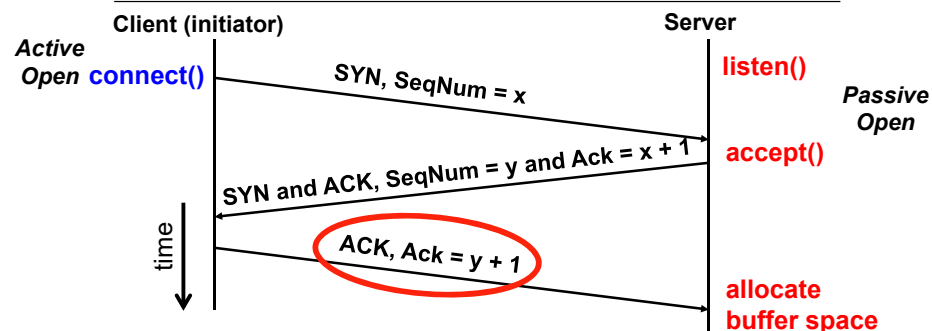


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.59

Denial of Service Vulnerability



- SYN attack: send a huge number of SYN messages
 - Causes victim to commit resources (768 byte TCP/IP data structure)
- Alternatives: Do not commit resources until receive final ACK
 - **SYN Cache**: when SYN received, put small entry into cache (using hash) and send SYN/ACK, If receive ACK, then put into listening socket
 - **SYN Cookie**: when SYN received, encode connection info into sequence number/other TCP header blocks, decode on ACK


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.60

Close Connection

- Goal: both sides agree to close the connection
- 4-way connection tear down



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.61

Use of TCP: Sockets


- **Socket**: an abstraction of a network I/O queue
 - Embodies one side of a communication channel
 - » Same interface regardless of location of other end
 - » Could be local machine (called “UNIX socket”) or remote machine (called “network socket”)
 - First introduced in 4.2 BSD UNIX: big innovation at time
- Using Sockets for Client-Server (C/C++ interface):
 - On server: set up “server-socket”
 - » Create socket, Bind to protocol (TCP), local address, port
 - » Call listen(): tells server socket to accept incoming requests
 - » Multiple accept() calls on socket to accept incoming connection requests
 - » Each successful accept() returns a new socket for a new connection
 - On client:
 - » Create socket, Bind to protocol (TCP), remote address, port
 - » Perform connect() on socket to make connection
 - » If connect() successful, have socket connected to server
- **Network Address Translation (NAT)**:
 - Local subnet (non-routable IP addresses) ⇒ external IP
 - Client-side firewall replaces local IP address/port combination with external IP address/new port
 - Firewall handles translation between different address domains using table of current connections

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.62

Recall: Socket Setup over TCP/IP



- Things to remember:
 - Connection involves 5 values:
[Client Addr, Client Port, Server Addr, Server Port, Protocol]
 - Often, Client Port “randomly” assigned
 - Server Port often “well known”
 - » 80 (web), 443 (secure web), 25 (sendmail), etc
 - » Well-known ports from 0—1023
- Network Address Translation (NAT) allows many internal connections (and/or hosts) with a single external IP address

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.63

Remote Procedure Call (RPC)

- Raw messaging is a bit too low-level for programming
 - Must wrap up information into message at source
 - Must decide what to do with message at destination
 - May need to sit and wait for multiple messages to arrive
- Another option: Remote Procedure Call (RPC)
 - Calls a procedure on a remote machine
 - Client calls:


```
remoteFileSystem→Read("rutabaga");
```
 - Translated automatically into call on server:


```
fileSys→Read("rutabaga");
```

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.64

RPC Implementation

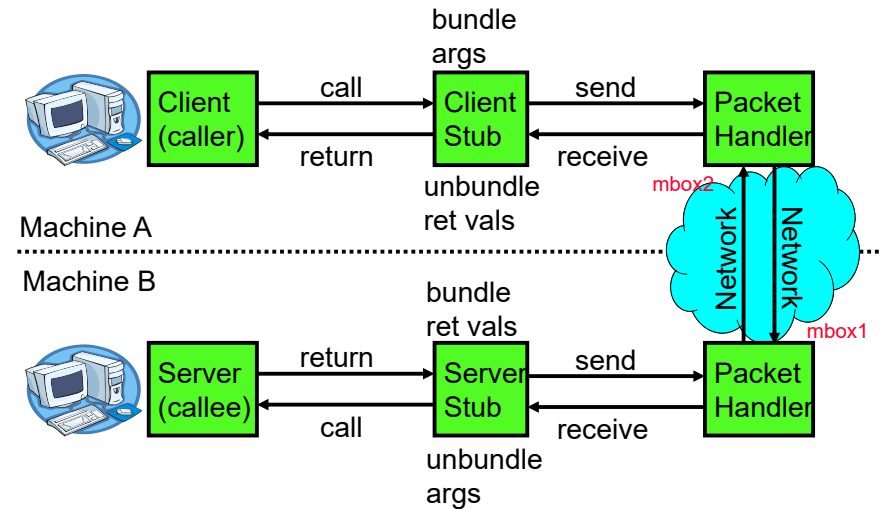
- Request-response message passing (under covers!)
- “Stub” provides glue on client/server
 - Client stub is responsible for “marshalling” arguments and “unmarshalling” the return values
 - Server-side stub is responsible for “unmarshalling” arguments and “marshalling” the return values.
- **Marshalling** involves (depending on system)
 - Converting values to a canonical form, serializing objects, copying arguments passed by reference, etc.

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.65

RPC Information Flow



4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.66

RPC Details (1/3)

- Equivalence with regular procedure call
 - Parameters \Leftrightarrow Request Message
 - Result \Leftrightarrow Reply message
 - Name of Procedure: Passed in request message
 - Return Address: mbox2 (client return mail box)
- Stub generator: Compiler that generates stubs
 - Input: interface definitions in an “interface definition language (IDL)”
 - » Contains, among other things, types of arguments/return
 - Output: stub code in the appropriate source language
 - » Code for client to pack message, send it off, wait for result, unpack result and return to caller
 - » Code for server to unpack message, call procedure, pack results, send them off

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.67

RPC Details (2/3)

- Cross-platform issues:
 - What if client/server machines are different architectures/languages?
 - » Convert everything to/from some canonical form
 - » Tag every item with an indication of how it is encoded (avoids unnecessary conversions)
- How does client know which mbox to send to?
 - Need to translate name of remote service into network endpoint (Remote machine, port, possibly other info)
 - **Binding**: the process of converting a user-visible name into a network endpoint
 - » This is another word for “naming” at network level
 - » Static: fixed at compile time
 - » Dynamic: performed at runtime

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.68

RPC Details (3/3)

- Dynamic Binding
 - Most RPC systems use dynamic binding via name service
 - » Name service provides dynamic translation of service → mbox
 - Why dynamic binding?
 - » Access control: check who is permitted to access service
 - » Fail-over: If server fails, use a different one
- What if there are multiple servers?
 - Could give flexibility at binding time
 - » Choose unloaded server for each new client
 - Could provide same mbox (router level redirect)
 - » Choose unloaded server for each new request
 - » Only works if no state carried from one call to next
- What if multiple clients?
 - Pass pointer to client-specific return mbox in request

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.69

Problems with RPC: Non-Atomic Failures

- Different failure modes in dist. system than on a single machine
- Consider many different types of failures
 - User-level bug causes address space to crash
 - Machine failure, kernel bug causes all processes on same machine to fail
 - Some machine is compromised by malicious party
- Before RPC: whole system would crash/die
- After RPC: One machine crashes/compromised while others keep working
- Can easily result in inconsistent view of the world
 - Did my cached data get written back or not?
 - Did server do what I requested or not?
- Answer? Distributed transactions/Byzantine Commit

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.70

Problems with RPC: Performance

- Cost of Procedure call « same-machine RPC « network RPC
- Means programmers must be aware that RPC is not free
 - Caching can help, but may make failure handling complex

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.71

Cross-Domain Communication/ Location Transparency

- How do address spaces communicate with one another?
 - Shared Memory with Semaphores, monitors, etc...
 - File System
 - Pipes (1-way communication)
 - “Remote” procedure call (2-way communication)
- RPC’s can be used to communicate between address spaces on different machines or the same machine
 - Services can be run wherever it’s most appropriate
 - Access to local and remote services looks the same
- Examples of RPC systems:
 - CORBA (Common Object Request Broker Architecture)
 - DCOM (Distributed COM)
 - RMI (Java Remote Method Invocation)


4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.72

Microkernel operating systems

- Example: split kernel into application-level servers.
 - File system looks remote, even though on same machine



- Why split the OS into separate domains?
 - Fault isolation: bugs are more isolated (build a firewall)
 - Enforces modularity: allows incremental upgrades of pieces of software (client or server)
 - Location transparent: service can be local or remote
 - » For example in the X windowing system: Each X client can be on a separate machine from X server; Neither has to run on the machine with the frame buffer.

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.73

Summary (1/2)

- Two-phase commit: distributed decision making
 - First, make sure everyone guarantees they will commit if asked (prepare)
 - Next, ask everyone to commit
- Byzantine General's Problem: distributed decision making with malicious failures
 - One general, $n-1$ lieutenants: some number of them may be malicious (often " f " of them)
 - All non-malicious lieutenants must come to same decision
 - If general not malicious, lieutenants must follow general
 - Only solvable if $n \geq 3f+1$
- Blockchain protocols
 - Cryptographically-driven ordering protocol
 - Could be used for distributed decision making

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.74

Summary (2/2)

- Internet Protocol (IP): Datagram packet delivery
 - Used to route messages through routes across globe
 - 32-bit addresses, 16-bit ports
- DNS: System for mapping from names \Rightarrow IP addresses
 - Hierarchical mapping from authoritative domains
 - Recent flaws discovered
- Ordered messages:
 - Use sequence numbers and reorder at destination
- Reliable messages:
 - Use Acknowledgements
- TCP: Reliable byte stream between two processes on different machines over Internet (read, write, flush)
 - Uses window-based acknowledgement protocol
 - Congestion-avoidance dynamically adapts sender window to account for congestion in network
- Remote Procedure Call (RPC): Call procedure on remote machine
 - Provides same interface as procedure
 - Automatic packing and unpacking of arguments without user programming (in stub)

4/21/20

Kubiatowicz CS162 © UCB Spring 2020

Lec 22.75